

```
!pip install torch torchvision matplotlib seaborn scikit-learn
```

```
Requirement already satisfied: torch in
/usr/local/lib/python3.11/dist-packages (2.5.1+cu121)
Requirement already satisfied: torchvision in
/usr/local/lib/python3.11/dist-packages (0.20.1+cu121)
Requirement already satisfied: matplotlib in
/usr/local/lib/python3.11/dist-packages (3.10.0)
Requirement already satisfied: seaborn in
/usr/local/lib/python3.11/dist-packages (0.13.2)
Requirement already satisfied: scikit-learn in
/usr/local/lib/python3.11/dist-packages (1.6.1)
Requirement already satisfied: filelock in
/usr/local/lib/python3.11/dist-packages (from torch) (3.17.0)
Requirement already satisfied: typing-extensions>=4.8.0 in
/usr/local/lib/python3.11/dist-packages (from torch) (4.12.2)
Requirement already satisfied: networkx in
/usr/local/lib/python3.11/dist-packages (from torch) (3.4.2)
Requirement already satisfied: jinja2 in
/usr/local/lib/python3.11/dist-packages (from torch) (3.1.5)
Requirement already satisfied: fsspec in
/usr/local/lib/python3.11/dist-packages (from torch) (2024.10.0)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.1.105 in
/usr/local/lib/python3.11/dist-packages (from torch) (12.1.105)
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.1.105
in /usr/local/lib/python3.11/dist-packages (from torch) (12.1.105)
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.1.105 in
/usr/local/lib/python3.11/dist-packages (from torch) (12.1.105)
Requirement already satisfied: nvidia-cudnn-cu12==9.1.0.70 in
/usr/local/lib/python3.11/dist-packages (from torch) (9.1.0.70)
Requirement already satisfied: nvidia-cublas-cu12==12.1.3.1 in
/usr/local/lib/python3.11/dist-packages (from torch) (12.1.3.1)
Requirement already satisfied: nvidia-cufft-cu12==11.0.2.54 in
/usr/local/lib/python3.11/dist-packages (from torch) (11.0.2.54)
Requirement already satisfied: nvidia-curand-cu12==10.3.2.106 in
/usr/local/lib/python3.11/dist-packages (from torch) (10.3.2.106)
Requirement already satisfied: nvidia-cusolver-cu12==11.4.5.107 in
/usr/local/lib/python3.11/dist-packages (from torch) (11.4.5.107)
Requirement already satisfied: nvidia-cuspars-cu12==12.1.0.106 in
/usr/local/lib/python3.11/dist-packages (from torch) (12.1.0.106)
Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in
/usr/local/lib/python3.11/dist-packages (from torch) (2.21.5)
Requirement already satisfied: nvidia-nvtx-cu12==12.1.105 in
/usr/local/lib/python3.11/dist-packages (from torch) (12.1.105)
Requirement already satisfied: triton==3.1.0 in
/usr/local/lib/python3.11/dist-packages (from torch) (3.1.0)
Requirement already satisfied: sympy==1.13.1 in
/usr/local/lib/python3.11/dist-packages (from torch) (1.13.1)
Requirement already satisfied: nvidia-nvjitlink-cu12 in
/usr/local/lib/python3.11/dist-packages (from nvidia-cusolver-
```

```
cul2==11.4.5.107->torch) (12.6.85)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in
/usr/local/lib/python3.11/dist-packages (from sympy==1.13.1->torch)
(1.3.0)
Requirement already satisfied: numpy in
/usr/local/lib/python3.11/dist-packages (from torchvision) (1.26.4)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in
/usr/local/lib/python3.11/dist-packages (from torchvision) (11.1.0)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (1.3.1)
Requirement already satisfied: cycler>=0.10 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (4.55.5)
Requirement already satisfied: kiwisolver>=1.3.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (1.4.8)
Requirement already satisfied: packaging>=20.0 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (24.2)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (3.2.1)
Requirement already satisfied: python-dateutil>=2.7 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (2.8.2)
Requirement already satisfied: pandas>=1.2 in
/usr/local/lib/python3.11/dist-packages (from seaborn) (2.2.2)
Requirement already satisfied: scipy>=1.6.0 in
/usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.13.1)
Requirement already satisfied: joblib>=1.2.0 in
/usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in
/usr/local/lib/python3.11/dist-packages (from scikit-learn) (3.5.0)
Requirement already satisfied: pytz>=2020.1 in
/usr/local/lib/python3.11/dist-packages (from pandas>=1.2->seaborn)
(2024.2)
Requirement already satisfied: tzdata>=2022.7 in
/usr/local/lib/python3.11/dist-packages (from pandas>=1.2->seaborn)
(2025.1)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.7-
>matplotlib) (1.17.0)
Requirement already satisfied: MarkupSafe>=2.0 in
/usr/local/lib/python3.11/dist-packages (from jinja2->torch) (3.0.2)
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
import torch
import torch.nn as nn
import torch.optim as optim
```

```

from torchvision import datasets, transforms, models
import os
from sklearn.metrics import confusion_matrix, classification_report
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from torch.optim import lr_scheduler
from torchvision.models import regnet_y_8gf # Latest RegNet variant

# Define data transformations for data augmentation and normalization
data_transforms = {
    'train': transforms.Compose([
        transforms.RandomResizedCrop(224),
        transforms.RandomHorizontalFlip(),
        transforms.ColorJitter(brightness=0.2, contrast=0.2,
saturation=0.2, hue=0.2),
        transforms.RandomRotation(30),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224,
0.225]) # ImageNet mean and std
    ]),
    'test': transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224,
0.225]) # ImageNet mean and std
    ]),
}

# Define the data directory
data_dir = '/content/drive/MyDrive/Images'

# Create data loaders
image_datasets = {x: datasets.ImageFolder(os.path.join(data_dir, x),
data_transforms[x]) for x in ['train', 'test']}
dataloaders = {x: torch.utils.data.DataLoader(image_datasets[x],
batch_size=16, shuffle=True, num_workers=4) for x in ['train',
'test']}
dataset_sizes = {x: len(image_datasets[x]) for x in ['train', 'test']}
print(dataset_sizes)

class_names = image_datasets['train'].classes

{'train': 3902, 'test': 977}

```

```

/usr/local/lib/python3.11/dist-packages/torch/utils/data/
dataloader.py:617: UserWarning: This DataLoader will create 4 worker
processes in total. Our suggested max number of worker in current
system is 2, which is smaller than what this DataLoader is going to

```

create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.

```
warnings.warn()
```

```
# Load the pre-trained RegNetY model (you can choose different variants of RegNet like 400MF, 8GF, etc.)
```

```
model = regnet_y_8gf(pretrained=True)
```

```
# Efficiently modify the final fully connected layer for 6 classes
```

```
model.fc = nn.Linear(model.fc.in_features, 6)
```

```
/usr/local/lib/python3.11/dist-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.
```

```
warnings.warn()
```

```
/usr/local/lib/python3.11/dist-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing `weights=RegNet_Y_8GF_Weights.IMAGENET1K_V1`. You can also use `weights=RegNet_Y_8GF_Weights.DEFAULT` to get the most up-to-date weights.
```

```
warnings.warn(msg)
```

```
Downloading: "https://download.pytorch.org/models/regnet_y_8gf-d0d0e4a8.pth" to /root/.cache/torch/hub/checkpoints/regnet_y_8gf-d0d0e4a8.pth
```

```
100%|██████████| 151M/151M [00:03<00:00, 49.7MB/s]
```

```
# Define the loss function and optimizer
```

```
criterion = nn.CrossEntropyLoss()
```

```
optimizer = optim.Adam(model.parameters(), lr=0.001) # Use Adam optimizer for better convergence
```

```
# Learning Rate Scheduler (use cosine annealing for RegNet's dynamic scaling)
```

```
scheduler = lr_scheduler.CosineAnnealingLR(optimizer, T_max=10)
```

```
# Move the model to the GPU if available
```

```
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

```
model = model.to(device)
```

```
# Training loop with dynamic scheduler and more efficient processing
```

```
num_epochs = 15
```

```
best_model_wts = model.state_dict()
```

```
best_acc = 0.0
```

```
train_losses, test_losses = [], []
```

```
train_accs, test_accs = [], []
```

```

for epoch in range(num_epochs):
    print(f'Epoch {epoch+1}/{num_epochs}')
    print('-' * 10)

    for phase in ['train', 'test']:
        if phase == 'train':
            model.train()
        else:
            model.eval()

        running_loss = 0.0
        running_corrects = 0

        for inputs, labels in dataloaders[phase]:
            inputs = inputs.to(device)
            labels = labels.to(device)

            optimizer.zero_grad()

            with torch.set_grad_enabled(phase == 'train'):
                outputs = model(inputs)
                _, preds = torch.max(outputs, 1)
                loss = criterion(outputs, labels)

                if phase == 'train':
                    loss.backward()
                    optimizer.step()

            running_loss += loss.item() * inputs.size(0)
            running_corrects += torch.sum(preds == labels.data)

        epoch_loss = running_loss / dataset_sizes[phase]
        epoch_acc = running_corrects.double() / dataset_sizes[phase]
        if phase == 'train':
            train_losses.append(epoch_loss)
            train_accs.append(epoch_acc)
        else:
            test_losses.append(epoch_loss)
            test_accs.append(epoch_acc)

    print(f'{phase} Loss: {epoch_loss:.4f} Acc: {epoch_acc:.4f}')

    # Deep copy the model if it achieves better performance
    if phase == 'test' and epoch_acc > best_acc:
        best_acc = epoch_acc
        best_model_wts = model.state_dict()

    scheduler.step() # Update learning rate using cosine annealing

```

```
# Load best model weights
model.load_state_dict(best_model_wts)
```

```
Epoch 1/15
```

```
-----
train Loss: 0.8761 Acc: 0.6991
test Loss: 0.7128 Acc: 0.7851
```

```
Epoch 2/15
```

```
-----
train Loss: 0.5385 Acc: 0.8175
test Loss: 0.3064 Acc: 0.8915
```

```
Epoch 3/15
```

```
-----
train Loss: 0.4485 Acc: 0.8432
test Loss: 0.1877 Acc: 0.9314
```

```
Epoch 4/15
```

```
-----
train Loss: 0.3551 Acc: 0.8842
test Loss: 0.1486 Acc: 0.9529
```

```
Epoch 5/15
```

```
-----
train Loss: 0.3039 Acc: 0.8949
test Loss: 0.1022 Acc: 0.9683
```

```
Epoch 6/15
```

```
-----
train Loss: 0.2500 Acc: 0.9188
test Loss: 0.0527 Acc: 0.9826
```

```
Epoch 7/15
```

```
-----
train Loss: 0.2000 Acc: 0.9341
test Loss: 0.0382 Acc: 0.9867
```

```
Epoch 8/15
```

```
-----
train Loss: 0.1709 Acc: 0.9441
test Loss: 0.0461 Acc: 0.9846
```

```
Epoch 9/15
```

```
-----
train Loss: 0.1333 Acc: 0.9564
test Loss: 0.0397 Acc: 0.9785
```

```
Epoch 10/15
```

```
-----
train Loss: 0.1092 Acc: 0.9634
test Loss: 0.0342 Acc: 0.9826
```

```
Epoch 11/15
```

```
-----
train Loss: 0.0915 Acc: 0.9669
test Loss: 0.0376 Acc: 0.9826
```

```
Epoch 12/15
```

```
-----
train Loss: 0.1070 Acc: 0.9654
```

```

test Loss: 0.0304 Acc: 0.9867
Epoch 13/15
-----
train Loss: 0.1126 Acc: 0.9618
test Loss: 0.0342 Acc: 0.9857
Epoch 14/15
-----
train Loss: 0.1325 Acc: 0.9526
test Loss: 0.0424 Acc: 0.9836
Epoch 15/15
-----
train Loss: 0.1871 Acc: 0.9400
test Loss: 0.0641 Acc: 0.9826

<All keys matched successfully>

# Save the trained model
torch.save(model.state_dict(), 'best_classification_model_regnet.pth')

# Evaluate model performance
def evaluate_model(model, dataloader):
    model.eval()
    all_preds = []
    all_labels = []

    with torch.no_grad():
        for inputs, labels in dataloaders['test']: # Using test set
            for evaluation
                inputs = inputs.to(device)
                labels = labels.to(device)
                outputs = model(inputs)
                _, preds = torch.max(outputs, 1)
                all_preds.extend(preds.cpu().numpy())
                all_labels.extend(labels.cpu().numpy())

    cm = confusion_matrix(all_labels, all_preds)
    cr = classification_report(all_labels, all_preds,
target_names=class_names)
    return cm, cr

# Get confusion matrix and classification report
cm, cr = evaluate_model(model, dataloaders['test'])
print("Classification Report:\n", cr)


```

Classification Report:

	precision	recall	f1-score	support
Aphids	0.99	0.96	0.98	164
Army worm	0.96	1.00	0.98	162
Bacterial blight	0.99	0.97	0.98	160
Healthy	0.98	1.00	0.99	165

Powdery mildew	0.99	1.00	1.00	166
Target spot	0.98	0.96	0.97	160
accuracy			0.98	977
macro avg	0.98	0.98	0.98	977
weighted avg	0.98	0.98	0.98	977

```
# Visualize Confusion Matrix
plt.figure(figsize=(10, 7))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=class_names, yticklabels=class_names)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```



```
from PIL import Image
# Example image prediction with probabilities
image_path = '/content/target.jpeg' # Replace with your test image path
image = Image.open(image_path)
preprocess = transforms.Compose([
```



```

        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225]))
    ])
    input_tensor = preprocess(image)
    input_batch = input_tensor.unsqueeze(0).to(device)

    # Perform inference
    model.eval()
    with torch.no_grad():
        output = model(input_batch)
        probabilities = torch.nn.functional.softmax(output, dim=1)
        confidence, predicted_class = torch.max(probabilities, 1)

    # Get predicted class and probability
    predicted_class_name = class_names[predicted_class.item()]
    predicted_prob = probabilities[0][predicted_class].item()

    print(f'The predicted class is: {predicted_class_name} with a
    confidence of {predicted_prob:.4f}')

    # Show the image with prediction
    plt.imshow(np.array(image))
    plt.axis('off')
    plt.text(10, 10, f'Predicted: {predicted_class_name}
    ({predicted_prob*100:.2f}%)', fontsize=12, color='white',
    backgroundColor='red')
    plt.show()

```

The predicted class is: Target spot with a confidence of 0.7808

Predicted: Target spot (78.08%)



```
# Ensure all tensors in train_accs and test_accs are converted to
Python floats
train_accs = [acc.cpu().item() if isinstance(acc, torch.Tensor) else
acc for acc in train_accs]
test_accs = [acc.cpu().item() if isinstance(acc, torch.Tensor) else
acc for acc in test_accs]

# Plotting
epochs = range(1, num_epochs + 1)
plt.figure(figsize=(12, 5))

# Loss plot
plt.subplot(1, 2, 1)
plt.plot(epochs, train_losses, label='Train Loss')
plt.plot(epochs, test_losses, label='Test Loss')
plt.title('Loss Curve')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

# Accuracy plot
plt.subplot(1, 2, 2)
plt.plot(epochs, train_accs, label='Train Accuracy')
plt.plot(epochs, test_accs, label='Test Accuracy')
plt.title('Accuracy Curve')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
```

```
plt.legend()  
plt.tight_layout()  
plt.show()
```

